



## ITTE105b – Analytics Application Model Deployment – Estimation of Obesity Levels Based on Eating Habits and Physical Condition

### Objectives

The objective of this activity is to deploy the best-performing machine learning model that predicts obesity levels based on individual health and lifestyle factors, and implement it as a web service using FastAPI and Docker.

#### I. Preparation of the Jupyter Notebook for Model Deployment

The notebook file [obesity levels](#) was used to train the model, with Support Vector Machines (SVM) identified as the best-performing model based on accuracy and other classification metrics. Before deploying the trained model, all necessary components, including the model itself as well as the preprocessing tools, such as the encoder and scaler, were serialized using the **joblib** library, which is designed for saving and loading Python objects.

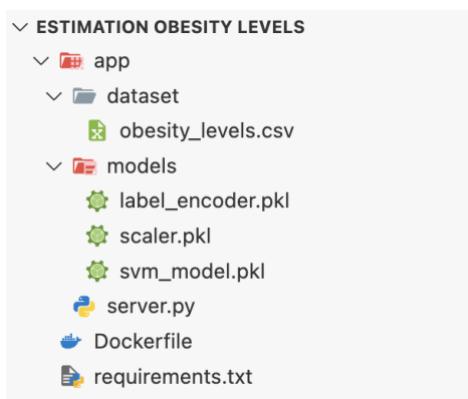
```
1. ...
2. import joblib;
3. ...
4.
5. svm = SVC(C=1000, gamma=0.001, probability=True)
6. svm.fit(X_train, y_train)
7.
8. joblib.dump(svm, 'svm_model.pkl')
9. joblib.dump(scaler, 'scaler.pkl')
10. joblib.dump(encoder, 'label_encoder.pkl')
```

#### II. Model Deployment using Fast API and Docker

After developing the model trained using SVM and serializing the preprocessing components, the next step is to deploy it as a web service using FastAPI and Docker. FastAPI is a Python framework used for building web APIs, while Docker is a platform that enables containerization of applications to ensure consistent deployment across different environments.

##### A. Folder Structure

The project folder [obesity\\_levels\\_pred](#) follows the structure shown below.





The **app directory** acts as the core component which contains the application logic, dataset and serialized model files necessary for the operation of Fast API service.

The **DockerFile** defines the instructions for building the Docker Image of the Fast API application. It specifies the base environment, dependencies, working directory, and startup commands.

```
1. FROM python:3.11
2.
3. WORKDIR /code
4.
5. COPY ./requirements.txt /code/requirements.txt
6.
7. RUN pip install --no-cache-dir -r /code/requirements.txt
8.
9. COPY ./app /code/app
10.
11. EXPOSE 8000
12.
13. CMD ["uvicorn", "app.server:app", "--host", "0.0.0.0", "--port", "8000"]
```

The **requirement.txt** file includes all Python dependencies required by the application, including **FastAPI**, **Scikit-learn**, **NumPy**, **Uvicorn** and **Pandas**.

```
1. scikit-learn==1.3.1
2. fastapi
3. numpy
4. uvicorn
5. pandas
```

The **app/dataset** folder contains the dataset used for training and validating the model. Meanwhile, the **app/models** folder includes the trained SVM model for predicting obesity levels, along with the serialized preprocessing tools generated from the **jupyter notebook file**. These tools include the **StandardScaler** object, which normalizes feature inputs during preprocessing, and the **LabelEncoder** object which converts categorical variables into numeric form.

The **server.py** implements the FastAPI web service. It loads the serialized model and preprocessing tools, defines API endpoints, and handles HTTP requests for prediction. The API receives input data, applies scaling and encoding, and returns predicted obesity levels and confidence scores in JSON format.

```
1. ....
2. def predict(data: dict):
3. ....
4.     # Scale the data from the dataframe
5.     new_data_scaled = scaler.transform(new_data)
6.     # Predict the obesity levels and probabilities
7.     prediction = svm.predict(new_data_scaled)
8.     probabilities = svm.predict_proba(new_data_scaled)
9.     # Get individual confidence scores for each predictions
10.    results = []
11.    for i in range(len(prediction)):
12.        pred_class = encoder.inverse_transform([prediction[i]])[0]
13.        confidence = float(probabilities[i][prediction[i]])
14.        results.append({
15.            "predicted_class": pred_class,
16.            "confidence": round(confidence, 2)
17.        })
18.
19.    return {"results": results}
```



## B. Build the Docker Image

To build the Docker image using Visual Studio Code (VSCode), first open your project folder in the editor. Once the folder is open, launch the terminal within VSCode and run the following **docker build** command. The image is labeled as **obesity\_levels\_pred\_img**. Once the command is executed successfully, Docker processes each instruction from the **Dockerfile**, installing dependencies, copying the FastAPI application files, and setting up the environment.

```
1. docker build -t obesity_levels_pred_img .
```

After successfully building the Docker image, the next step is to run the container to make the FastAPI service accessible.

```
1. docker run --name obesity_levels_pred -p 8000:8000 obesity_levels_pred_img
```

When executed, Docker launches the containerized FastAPI application and exposes the API endpoints defined in the **server.py** file. If the application runs successfully without errors, the corresponding container will appear in the Docker dashboard, indicating that the deployment is active.

	Name	Container ID	Image	Port(s)	Actions
	obesity_levels_pred	3bc9e8164422	obesity_levels_pred	8000:8000	

The root endpoint can be accessed at **localhost:8000**, which can also be used to verify whether the API is running.



```
localhost:8000
Pretty-print □
{"message": "Model API"}
```

## C. Access the Fast API Document

Once the Docker container is successfully running, the deployed model can be accessed through **Swagger UI** or **Redoc**, both of which are automatically generated interactive documentation provided by FastAPI. These interfaces allow users to explore, interact with, and test the API endpoints directly from a web browser.

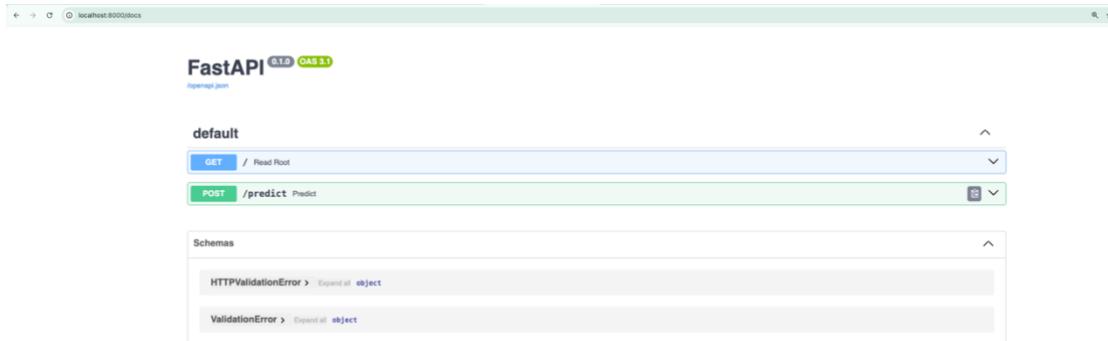
## D. Steps in accessing the interactive documentations:

1. Open a web browser and go to <http://localhost:8000/docs>. This opens the **Swagger UI**, an interactive documentation tool that displays all available endpoints of the API. The API contains the following end points.
  - The root endpoint / provides a simple connection test message.



pioneering excellence

- The /predict endpoint allows users to send input data and receive prediction results.



A screenshot of a web browser displaying the FastAPI documentation. The title bar shows "FastAPI 0.1.0 (GAS 2.1) openapi.json". Below the title, there's a "default" section with two buttons: "GET / Read Root" and "POST /predict Predict". Underneath these buttons is a "Schemas" section containing two items: "HTTPValidationError > Expand all object" and "ValidationError > Expand all object".

- After accessing the FastAPI documentation, you can test the deployed model's predictions by providing sample input data in the required format through the **/predict** endpoint.



A screenshot of the FastAPI /predict endpoint documentation. The top part shows the endpoint details: "POST /predict Predict". Below it, a description states: "Predicts obesity levels based on individual health and lifestyle features. Args: data (dict): Input dictionary containing "features" key with a list of feature arrays. Example: { "features": [ [85, 2, 1, 25, 2.0, 1, 1.70, 1, 3, 0], [68, 1, 0, 20, 3.0, 0, 1.65, 0, 2, 0] ] } Returns: dict: A dictionary containing predicted classes and confidence scores". The "Parameters" section indicates "No parameters". The "Request body" section is marked as "required" and set to "application/json". Below this, a "Request Body" section shows a JSON schema example:

```
{  
    "features": [  
        [85, 2, 1, 25, 2.0, 1, 1.70, 1, 3, 0],  
        [68, 1, 0, 20, 3.0, 0, 1.65, 0, 2, 0]  
    ]  
}
```

The sample input data follow the order of the input features defined in the developed model, as described in the following list of input feature descriptions.

Features	Description	Example Value
Weight (kg)	The individual's body weight in kilograms.	70
CAEC (Consumption of Food Between Meals)	Encoded frequency of eating between meals: 0 = No, 1 = Sometimes, 2 = Frequently, 3 = Always.	2
family_history_with_overweight	Indicates if the individual has family members with overweight or obesity: 0 = No, 1 = Yes.	1
Age	The individual's age in years.	25
CH2O (Water Intake per Day)	Average daily water consumption: 1 = Less than 1L, 2 = 1-2L, 3 = More than 2L.	2.0
Height (m)	The individual's height in meters.	1.70



*pioneering excellence*

Gender	Gender of the individual: 0 = Female, 1 = Male.	1
FCVC (Frequency of Vegetable Consumption)	Frequency of eating vegetables: 1 = Never, 2 = Sometimes, 3 = Always.	3
SMOKE	Indicates whether the individual smokes: 0 = No, 1 = Yes.	0

After entering the sample input data, click the “Execute” button to review the response section below. This section displays the predicted output along with the confidence score for each returned prediction.

Request URL  
`http://localhost:8000/predict`

Server response

Code	Details
200	<p>Response body</p> <pre>{ "results": [ { "predicted_class": "Overweight_Level_II", "confidence": 0.95 }, { "predicted_class": "Normal_Weight", "confidence": 0.57 } ] }</pre> <p>Download</p> <p>Response headers</p> <pre>access-control-allow-credentials: true access-control-allow-origin: http://localhost:8000 content-type: application/json date: Mon, 03 Nov 2025 01:07:12 GMT server: unicorn vary: Origin</pre>

Responses

Code	Description	Links
200	Successful Response	No links

### III. Integrate with Client Application

After verifying that the FastAPI model endpoint functions correctly, the next step is to integrate it with a web or mobile application interface which allows users to interact with the model through an intuitive front-end form instead of directly using Swagger UI. A web or mobile application framework can be used to send requests to the web service and process the responses and display the model's predictions to the user in a user-friendly manner.